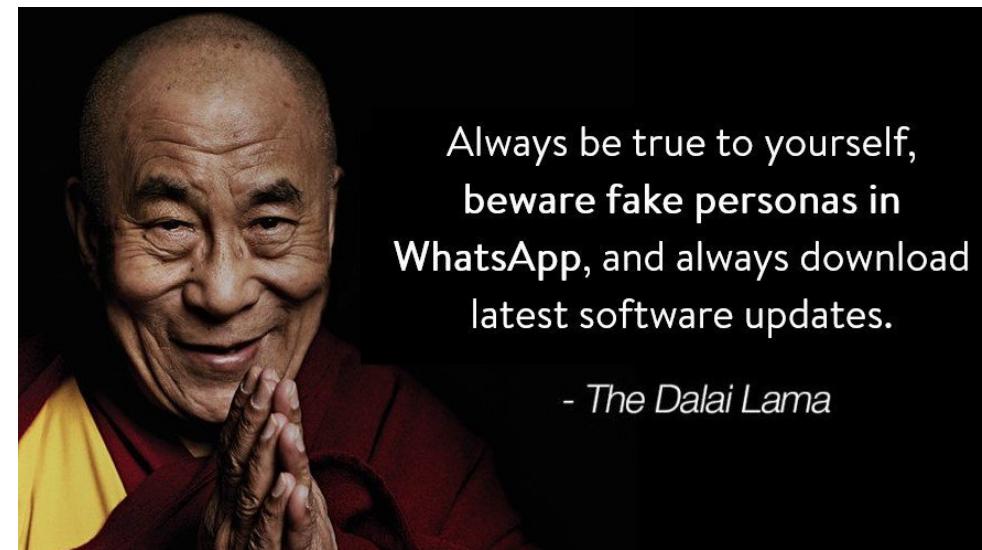


# Crypto 4: Public Key



Always be true to yourself,  
beware fake personas in  
WhatsApp, and always download  
latest software updates.

- *The Dalai Lama*

# Twitter Fight Last Year: Nick Vs Rust Rand\_Core Random Number Generators

- Rust (well, the 3rd party library for it) has an interface for "secure" Random Number Generators... But they aren't actually secure!
- EG, "ChaCha8Rng"
  - A *reduced round* stream cipher!
  - That has no update() function: no way of adding in entropy after seeding
  - And seed() takes only 32B total (no combining entropy!)
  - Oh, and no rollback resistance either
- **NONE** of the "Secure" RNGs are actually cryptographically secure...
  - Because none accept and consume arbitrarily long seeds or have an update to mix in more entropy
- When I say ONLY use HMAC\_DRBG, I mean it!
  - Use /dev/urandom and everything else you can think of to shove into HMAC\_DRBG

# And Vuln of the Day: CVE-2019-16303

- If you wrote an app in JHipster last year or before...
  - You probably want a password reset function...
- Password reset generates "random" URLs
  - But of course, they used a bad RNG!
- So generate a password request for your account
  - You get the RNGs state in the reset URL
- Now you can generate more password resets...
  - And predict what the "random" URL is...  
and take over any account you want!

# Reminder Of Our Primitives So-Far: Block Cipher

- Block Cipher: Takes a fixed sized message and fixed-sized key
  - $E(M, K)$ ,  $E_k(M)$
  - Corresponding inverse/decryption function  $D_k(M)$
  - Keyed permutation on an  $N$  bit block:  
If you don't know the key, it should be indistinguishable from a random permutation
  - If you change a single bit of either the input or the key, the output should look totally different
  - E.g. AES: 128b data blocks, keys are 128, 192, 256 (AES-128, AES-192, AES-256)
- Block Cipher Mode
  - A way of repeatedly applying a block cipher on a longer message:  
Goal is to make it independent under chosen plaintext attacks

# Reminder Of Primitives So-Far: Hash Function

- Hash takes an arbitrary message  $M$  and reduces it to a fixed size
  - Should be indistinguishable from a random number
  - Change a single bit on the input -> Output looks like a completely different random number
  - SHA-256, SHA-384, SHA-512: SHA2 family outputting 256b, 384b, 512b
  - SHA3-256, SHA3-384, SHA3-512: SHA3 family
- Irreversible & resists collisions
  - Intractable given  $H(X)$  to determine  $X$   
(1st Preimage Resistant)
  - Intractable given  $X$ ,  $H(X)$ , find  $X' \neq X$  such that  $H(X) = H(X')$   
(2nd Preimage Resistant)
  - Intractable to find any  $X$ ,  $X'$ ,  $X' \neq X$  such that  $H(X) = H(X')$   
(Collision Resistant)

# Reminder Of Primitives So-Far: MAC

- MAC takes an arbitrary message  $M$  and a key  $K$  creating a fixed-length tag
  - $\text{MAC}(M, K) \rightarrow T$
  - Without  $K$ , it is infeasible to create  $M'$  such that  $\text{MAC}(M', K) \rightarrow T$
  - Without  $K$ , it is infeasible to create  $M'$ ,  $T'$  such that  $\text{MAC}(M', K) \rightarrow T'$
  - But with  $K$ , of course you can create a valid  $M'$ ,  $T'$  pair
    - And for some MACs create  $M'$  which MACs to  $T$
- Several alternatives but only One True MAC to use:  
HMAC
  - Construct using hash functions to create a MAC:  
Has all the previous properties of a hash plus all the properties of a MAC

# Reminder Of Primitives So-Far: pRNG (Pseudo Random Number Generator)

- Three operations:
  - `seed(entropy)`: Set internal state based on arbitrarily long, truly random inputs
  - `update(entropy)`: Add in additional entropy  
Update with 0-entropy should not degrade internal state
  - `generate(length)`: Generate an  $n$  bit string that should be indistinguishable from random
- If you know the internal state it is fully predictable
- If you don't it should be indistinguishable from random
- HMAC\_DRBG is the absolute best
  - Also has rollback resistance, if you learned the internal state at time T, you can't predict previous outputs

# Public Key...

- All our previous primitives required a "miracle":
  - We somehow have to have Alice and Bob get a shared  $k$ .
- Enter Public Key cryptography: the miracle of modern cryptography
  - How starting Friday, but ***what*** today
- Three primitives:
  - Public Key Agreement (previous Ephemeral Diffie/Hellman)
  - Public Key Encryption
  - Public Key Signatures
- Based on some families of magic math...
  - For us, we will use some group-theory based primitives

# Public Key Agreement

- Alice and Bob have a channel...
  - There may be an eavesdropper ***but not a manipulator***
- The goal: Alice & Bob agree on a ***random*** value
  - This will be ***k*** for all subsequent communication
- When done, the key is thrown away
  - Designed to prevent an attacker who later recovers Alice or Bob's long lived secrets from finding ***k***.

# Reminder of Primitives So Far: Ephemeral Diffie/Hellman Key Exchange

- Public values: prime  $p$ , generator  $g$ 
  - Elliptic curve: different magic math, fewer bits (256b/384b instead of 2048b/3096b for the same security)
- Alice creates random  $a$ ,  $0 < a < p$ , computes  $A = g^a \text{ mod } p$ , sends it
- Bob creates random  $b$ ,  $0 < b < p$ , computers  $B = g^b \text{ mod } p$ , sends it
- Alice computes  $B^a \text{ mod } P = g^{ab} \text{ mod } P = K$
- Bob computes  $A^b \text{ mod } P = g^{ab} \text{ mod } P = K$
- Thought to be hard to go backwards (discrete log) to  $a$  given  $A$

# Public Key Encryption

- Alice has **two** keys:
  - $K_{pub}$ : Her public key, anyone can know
  - $K_{priv}$ : Her private key, a deep dark secret
    - Sometimes written as  $K_{alice}$ ,  $K^{-1}_{alice}$
- Anyone has access to Alice's public key
- For anyone to send a message to Alice:
  - Create a random session key  $k$ 
    - Used to encrypt the rest of the message
  - Encrypt  $k$  using Alice's  $K_{pub}$ .
- Only Alice can **decrypt** the message
  - The decryption function only works with  $K_{priv}$ !

# Public Key Signatures

- Once again, Alice has **two** keys:
  - $K_{pub}$ : Her public key, anyone can know
  - $K_{priv}$ : Her private key, a deep dark secret
- She can sign a message
  - Calculate  $H(M)$
  - $S(K_{priv}, H(M))$ : Sign  $H(M)$  with  $K_{priv}$ .
- Anyone can now verify
  - Recalculate  $H(M)$
  - $V(K_{pub}, S(K_{priv}, H(M)), H(M))$ : Verify that the signature was created with  $K_{priv}$

# Things To Remember...

- Public key is ***slow!***
  - Orders of magnitude slower than symmetric key
- Public key is based on delicate magic math
  - Discrete log in a group is the most common
  - RSA
  - Some new "post-quantum" magic...
- Some systems in particular are easy to get wrong
  - We will get to some of the epic crypto-fails later

# Our Roadmap For Public Key...

- Public Key:
  - Something **everyone** can know
- Private Key:
  - The secret belonging to a specific person
- Diffie/Hellman:
  - Provides key exchange with no pre-shared secret
- ElGamal & RSA:
  - Provide a message to a recipient only knowing the recipient's **public key**
- DSA & RSA signatures:
  - Provide a message that anyone can prove was generated with a **private key**

# Public Key Cryptography #1: RSA

- Alice generates two **large** primes, **p** and **q**
  - They should be generated randomly:  
Generate a large random number and then use a "primality test":  
A **probabilistic** algorithm that checks if the number is prime
- Alice then computes **n = p\*q** and  **$\phi(n) = (p-1)(q-1)$** 
  - $\phi(n)$  is Euler's totient function, in this case for a composite of two primes
  - **n** is big: 2048b to 4096b long!
- Choose random  **$2 < e < \phi(n)$** 
  - **e** also needs to be relatively prime to  $\phi(n)$  but it can be small
- Solve for  **$d = e^{-1} \bmod \phi(n)$** 
  - You can't solve for **d** without knowing  $\phi(n)$ , which requires knowing **p** and **q**
- **n, e** are public, **d, p, q, and  $\phi(n)$**  are secret

# RSA Encryption

- Bob can easily send a message  $m$  to Alice:
  - Bob computes  $c = m^e \text{ mod } n$
  - Without knowing  $d$ , it is believed to be intractable to compute  $m$  given  $c$ ,  $e$ , and  $n$
  - But if you can get  $p$  and  $q$ , you can get  $d$ :  
It is ***not known*** if there is a way to compute  $d$  without also being able to factor  $n$ , but it is known that if you can factor  $n$ , you can get  $d$ .
  - And factoring is ***believed*** to be hard to do
- Alice computes  $m = c^d \text{ mod } n = m^{ed} \text{ mod } n$
- Time for some math magic...

# RSA Encryption/Decryption, con't

- So we have:  $D(C, K_D) = (M^{e \cdot d}) \bmod n$
- Now recall that  $d$  is the multiplicative inverse of  $e$ , modulo  $\Phi(n)$ , and thus:

$$e \cdot d = 1 \bmod \Phi(n) \quad (\text{by definition})$$

$$e \cdot d - 1 = k \cdot \Phi(n) \quad \text{for some } k$$

- Therefore  $D(C, K_D) = M^{e \cdot d} \bmod n = (M^{e \cdot d - 1}) \cdot M \bmod n$   
 $= (M^{k\Phi(n)}) \cdot M \bmod n$   
 $= [(M^{\Phi(n)})^k] \cdot M \bmod n$   
 $= (1^k) \cdot M \bmod n \quad \text{by Euler's Theorem: } a^{\Phi(n)} \bmod n = 1$   
 $= M \bmod n = M$

(believed) Eve can recover  $M$  from  $C$  iff Eve can factor  $n=p \cdot q$

# But It Is Not That Simple...

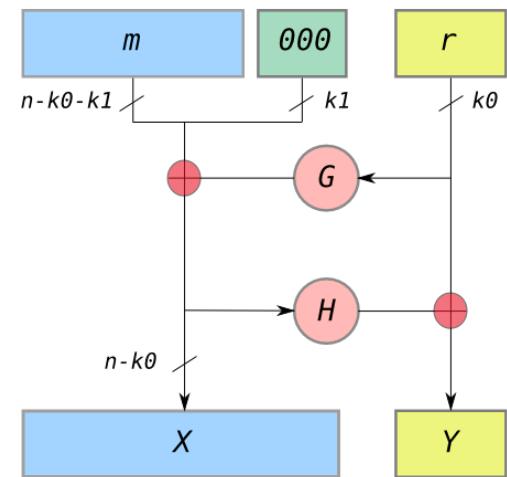
- What if Bob wants to send the same message to Alice twice?
  - Sends  $m^{e_a} \bmod n_a$  and then  $m^{e_a} \bmod n_a$
  - Oops, not IND-CPA!
- What if Bob wants to send a message to Alice, Carol, and Dave:
  - $m^{e_a} \bmod n_a$
  - $m^{e_b} \bmod n_b$
  - $m^{e_c} \bmod n_c$
  - This ends up leaking information an eavesdropper can use **especially** if  $3 = e_a = e_b = e_c$  !
- Oh, and problems if both **e** and **m** are small...
- As a result, you **can not** just use plain RSA:
  - You need to use a "padding" scheme that makes the input random but reversible



# RSA-OAEP

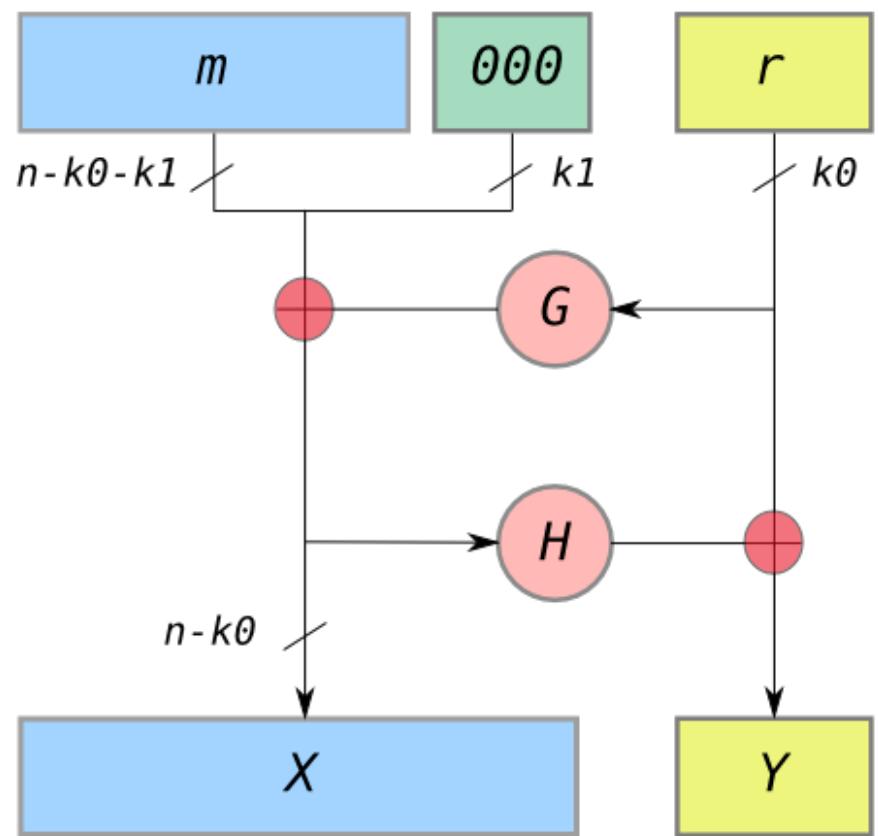
## (Optimal asymmetric encryption padding)

- A way of processing  $m$  with a hash function & random bits
  - Effectively "encrypts"  $m$  replacing it with  $X = [m, 0\dots] \oplus G(r)$ 
    - $G$  and  $H$  are hash functions (EG SHA-256)  
 $k_0 = \# \text{ of bits of randomness}, \text{len}(m) + k_1 + k_0 = n$
    - Then replaces  $r$  with  $Y = H(G(r) \oplus [m, 0\dots]) \oplus R$
    - This structure is called a "Feistel network":
      - It is always designed to be reversible.  
Many block ciphers are based on this concept applied multiple times with  $G$  and  $H$  being functions of  $k$  rather than just fixed operations
  - This is more than just block-cipher padding (which involves just adding simple patterns)
    - Instead it serves to both pad the bits and make the data to be encrypted "random"



# So How Does This Work?

- G and H are not (necessarily) reversible
  - EG, for OAEP it is a hash function:  
Designed to mix in the randomness and make it uniform
  - Needed for RSA because we want to only ever encrypt "random" values with the public key
  - And since  $r$  is random and  $G$  is a hash,  $m$  is xor'ed with random...
    - Which is then hashed and XOR'ed back into  $r$  to produce  $Y$
- But XOR is!
  - So we do  $H(X)$  xor  $Y$  to recover  $r$
  - And now  $G(r)$  xor  $X$  to recover  $m$



# But Its Not That Simple...

## Timing Attacks

Computer Science 161 Fall 2020

Weaver

- Using normal math, the **time** it takes for Alice to decrypt **c** depends on **c** and **d**
  - Ruh roh, this can leak information...
  - More complex RSA implementations take advantage of knowing **p** and **q** directly... but also leak timing
- People have used this to guess and then check the bits of **q** on OpenSSL
  - <http://crypto.stanford.edu/~dabo/papers/ssl-timing.pdf>
  - And even more subtle things are possible...

```
x = C
for j = 1 to n
    x = mod(x2, N)
    if dj == 1 then
        x = mod(xC, N)
    end if
next j
return x
```



# So How to Find Bob's Key?

- Lots of stuff later, but for now...

## ***The Leap of Faith!***

- Alice wants to talk to Bob:
  - "Hey, Bob, tell me your public key!"
- Now on all subsequent times...
  - "Hey, Bob, tell me your public key", and check to see if it is different from what Alice remembers
- Works assuming the ***first time*** Alice talks to Bob there isn't a Man-in-the-Middle
  - ssh uses this

# RSA Signatures...

- Alice computes a hash of the message  $H(m)$ 
  - Alice then computes  $s = (H(m))^d \bmod n$
- Anyone can then verify
  - $v = s^e \bmod m = ((H(m))^d)^e \bmod n = H(m)$
- Once again, there are "F-U"s...
  - Have to use a proper encoding scheme to do this properly and all sort of other traps
  - One particular trap: a scenario where the attacker can get Alice to repeatedly sign things (an "oracle")



# But Signatures Are Super Valuable...

- They are how we can prevent a MitM!
- If Bob knows Alice's key, and Alice knows Bob's...
- Alice doesn't just send a message to Bob...
  - But creates a random key  $k$ ...
  - Sends  $E(M, K_{sess})$ ,  $E(K_{sess}, B_{pub})$ ,  $S(H(M), A_{priv})$
- Only Bob can decrypt the message, and Bob can verify the message came from Alice
  - So Mallory is SOL!

# RSA Isn't The Only Public Key Algorithm

- Isn't RSA enough?
  - RSA isn't particularly compact or efficient: dealing with 2000b (comfortably secure) or 3000b (NSA-paranoia) bit operations
  - Can we get away with fewer bits?
    - Well, Diffie-Hellman isn't any better...
    - But ***elliptic curve*** Diffie-Hellman is
- RSA also had some patent issues
  - So an attempt to build public key algorithms around the Diffie-Hellman problem

# El-Gamal

- Just like Diffie-Hellman...
  - Select  $p$  and  $g$ 
    - These are public and can be shared:  
Note, they need to be carefully considered how to create  $p$  and  $g$ ...  
Math beyond the level of this class
  - Alice chooses  $x$  randomly as her private key
    - And publishes  $h = g^x \text{ mod } p$  as her public key
  - Bob, to encrypt  $m$  to Alice...
    - Selects a *random*  $y$ , calculates  $c_1 = g^y \text{ mod } p$ ,  $s = h^y \text{ mod } p = g^{xy} \text{ mod } p$ 
      - $s$  becomes a shared secret between Alice and Bob
    - Maps message  $m$  to create  $m'$ , calculates  $c_2 = m' * s \text{ mod } p$
  - Bob then sends  $\{c_1, c_2\}$

# El-Gamal Decryption

- Alice first calculates  $s = c_1^x \bmod p$ 
  - Then Alice calculates  $m' = c_2 * s^{-1} \bmod p$
  - Then Alice calculates the inverse of the mapping to get  $m$
- Of course, there are problems...
  - Attacker can always change  $m'$  to  $2m'$
  - What if Bob screws up and reuses  $y$ ?
  - $c_2 = m_1' * s \bmod p$   
 $c_2' = m_2' * s \bmod p$
  - Ruh roh, this leaks information:  
 $c_2 / c_2' = m_1' / m_2'$
  - So if you know  $m_1$ ...



# In Practice: Session Keys...

- You use the public key algorithm to encrypt/agree on a session key..
  - And then encrypt the real message with the session key
  - You **never** actually encrypt the message itself with the public key algorithm
  - Often a set of keys: encrypt and MAC keys that are separate in each direction
- Why?
  - Public key is **slow**... Orders of magnitude slower than symmetric key
  - Public key may cause weird effects:
    - EG, El Gamal where an attacker can change the message to  **$2m$** ...
      - If  $m$  had meaning, this would be a problem
      - But if it just changes the encryption and MAC keys, the main message won't decrypt

# DSA Signatures...

- Again, based on Diffie-Hellman
  - Two initial parameters, **L** and **N**, and a hash function **H**
  - **L** == key length, eg 2048
  - **N** <= **len(H)**, e.g. 256
  - An N-bit prime **q**, an L-bit prime **p** such that **p - 1** is a multiple of **q**, and **g = h<sup>(p-1)/q</sup> mod p** for some arbitrary **h** ( $1 < h < p - 1$ )
  - $\{p, q, g\}$  are public parameters
- Alice creates her own random private key **x < q**
- Public key **y = g<sup>x</sup> mod p**

# Alice's Signature...

- Create a random value  $k < q$ 
  - Calculate  $r = (g^k \bmod p) \bmod q$ 
    - If  $r = 0$ , start again
  - Calculate  $s = k^{-1} (H(m) + xr) \bmod q$ 
    - If  $s = 0$ , start again
  - Signature is  $\{r, s\}$  (Advantage over an El-Gamal signature variation: Smaller signatures)
- Verification
  - $w = s^{-1} \bmod q$
  - $u_1 = H(m) * w \bmod q$
  - $u_2 = r * w \bmod q$
  - $v = (g^{u_1} y^{u_2} \bmod p) \bmod q$
  - Validate that  $v = r$

# But Easy To Screw Up...

- **k** is not just a nonce... It must be random and **secret**
  - If you know **k**, you can calculate **x**
  - And even if you just reuse a random **k**...  
for two signatures **s<sub>a</sub>** and **s<sub>b</sub>**
    - A bit of algebra proves that  $k = (H_A - H_B) / (s_a - s_b)$
  - A good reference:
    - How knowing **k** tells you **x**:  
<https://rdist.root.org/2009/05/17/the-debian-pgp-disaster-that-almost-was/>
    - How two signatures tells you **k**:  
<https://rdist.root.org/2010/11/19/dsa-requirements-for-random-k-value/>



# And ***NOT*** theoretical: Sony Playstation 3 DRM

- The PS3 was designed to only run ***signed*** code
  - They used ECDSA as the signature algorithm
  - This prevents unauthorized code from running
  - They had an ***option*** to run alternate operating systems (Linux) that they then removed
- Of course this was catnip to reverse engineers
  - Best way to get people interested:  
***remove*** Linux from a device...
- It turns out one of the key authentication keys used to sign the firmware...
  - Ended up reusing the same k for multiple signatures!



# And **NOT** Theoretical: Android RNG Bug + Bitcoin

- OS Vulnerability in 2013  
Android "SecureRandom" wasn't actually secure!
  - Not only was it low entropy, it would occasionally return the **same value multiple times**
- Multiple Bitcoin wallet apps on Android were affected
  - "Pay B Bitcoin to Bob" is signed by Alice's public key using ECDSA
    - Message is broadcast publicly for all to see
    - So you'd have cases where "Pay B to Bob" and "Pay C to Carol" were signed with the same k
  - **So of course** someone scanned for **all** such Bitcoin transactions



# And *Still* Happens! Chromebook

- Chromebooks have a built in U2F "Security key"
  - Enables signatures using 256b ECDSA to validate to particular websites
- There was a bug in the secure hardware!
  - Instead of using a random  $k$  that was 256b long, a bug caused it to be 32b long!
  - So an attacker who had a signature could simply try all possible  $k$  values!
- Fortunately in this case the damage was slight: this is for authenticating to a single website: each site used its own private key
- But still...
- <https://www.chromium.org/chromium-os/u2f-ecdsa-vulnerability>



# So What To Use?

- Paranoids like me:  
Good libraries and use the parameters from NSA's CNSA suite
  - Open algorithms approved for Top Secret communication
  - Better yet, libraries that implement full protocols that use these under the hood!
- Symmetric cipher: AES: 256b
  - CFB mode, thankyouverymuch. Counter mode and modes which include counter mode can DIAF...
- Hash function: SHA-384
  - Use HMAC for MAC
- RSA: 3072b
- Diffie/Hellman: 3072b
- ECDH/ECDSA: P-384
- But really, this is extra paranoid:  
2048b RSA/DH, 256b EC, 128b AES, SHA-256 excellent in practice

# How Can We Communicate With Someone New?

- Public-key crypto gives us amazing capabilities to achieve confidentiality, integrity & authentication without shared secrets ...
- But how do we solve MITM attacks?
- How can we trust we have the true public key for someone we want to communicate with?
- Ideas?

# Trusted Authorities

- Suppose there's a party that everyone agrees to trust to confirm each individual's public key
  - Say the Governor of California
  - Issues with this approach?
    - How can everyone agree to trust them?
    - Scaling: huge amount of work; single point of failure ...
      - ... and thus Denial-of-Service concerns
    - How do you know you're talking to the right authority??



# Trust Anchors

- Suppose the trusted party distributes their key so everyone has it ...







# Trust Anchors

- Suppose the trusted party distributes their key so everyone has it ...
- We can then use this to bootstrap trust
  - As long as we have confidence in the decisions that that party makes

# Digital Certificates

- Certificate (“cert”) = signed claim about someone’s public key
  - More broadly: a signed *attestation* about some claim
- Notation:
  - { M }<sub>K</sub> = “message M encrypted with public key k”
  - { M }<sub>K<sup>-1</sup></sub> = “message M signed w/ private key for K”
- E.g. M = “Nick's public key is K<sub>Nick</sub> = **0xF32A99B...**”  
Cert: M,  
{“Nick's public key ... **0xF32A99B...**” }<sub>K<sup>-1</sup> Gavin</sub>  
**= 0x923AB95E12...9772F**

# Certificate



Gavin Newsom hereby asserts:

Nick's public key is  $K_{Nick} = 0xF32A99B\dots$

The signature for this statement using

$K^{-1}_{Gavin}$  is  $0x923AB95E12\dots9772F$

# Certificate



Gavin Newsom hereby asserts:

Nick's public key is  $K_{Nick} = 0xF32A99B\dots$

The signature for this statement using

$K^{-1}$  This is  $0x923AB95E12\dots 9772F$

# Certificate



Gavin Newsom hereby asserts:

Nick's public key is  $K_{Nick} = 0xF32A99B...$

The signature  $f$  is computed over all of this

$K^{-1}_{Gavin}$  is  $0x923AB95E12\ldots9772F$

# Certificate



Gavin Newsom hereby asserts:

Nick's public key is  $K_{Nick} = 0xF32A99B\dots$

The signature for this statement using

$K^{-1}_{Gavin}$  is  $0x923AB95E12\dots 9772F$

and can be  
*validated* using:

# Certificate



This:

Gavin Newsom hereby asserts:

Nick's public key is  $K_{Nick} =$

The signature for this state

$K^{-1}_{Gavin}$  is **0x923AB95**



# If We Find This Cert Shoved Under Our Door ...

- What can we figure out?
  - If we know Gavin's key, then whether he indeed signed the statement
  - If we trust Gavin's decisions, then we have confidence we really have Nick's key
- Trust = ?
  - Gavin won't willy-nilly sign such statements
  - Gavin won't let his private key be stolen

# Analyzing Certs Shoved Under Doors ...

- **How** we get the cert doesn't affect its utility
- **Who** gives us the cert doesn't matter
  - They're not any more or less trustworthy because they did
  - Possessing a cert doesn't establish any identity!
- However: if someone demonstrates they can decrypt data encrypted with  $K_{Nick}$ , then we have high confidence they possess  $K^{-1}_{Nick}$ 
  - Same for if they show they can sign “using”  $K^{-1}_{Nick}$

# Scaling Digital Certificates

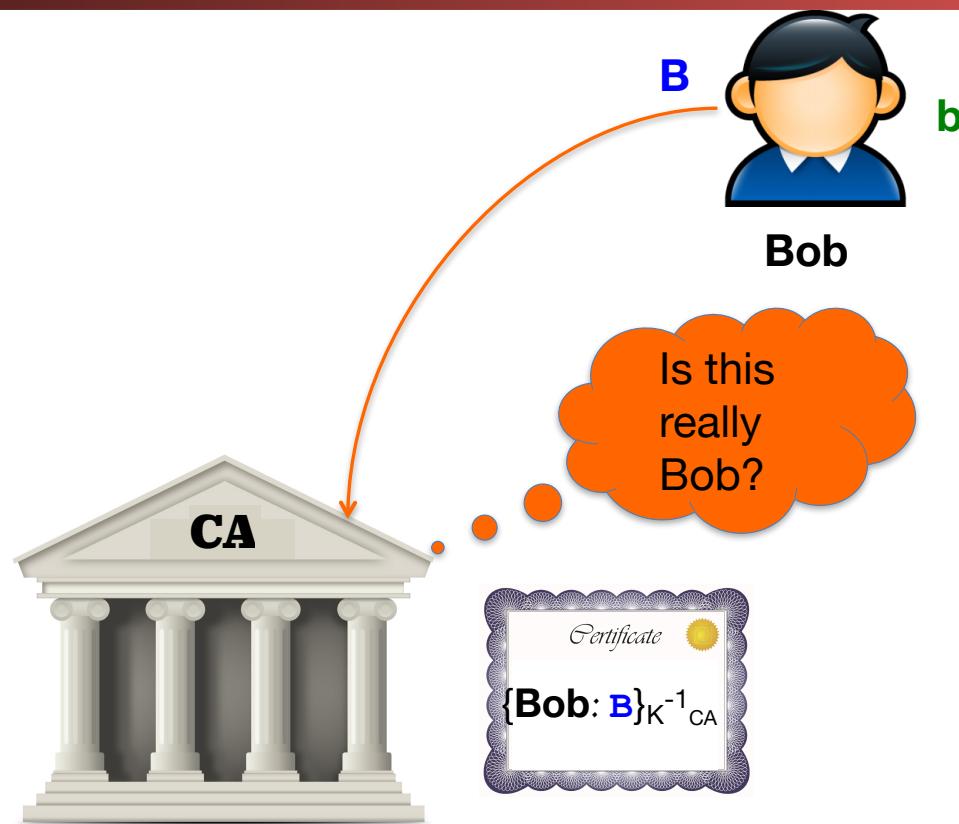
- How can this possibly scale? Surely Gavin can't sign everyone's public key!
- Approach #1: Introduce hierarchy via delegation
  - { "Michael V. Drake's public key is 0x... and I trust him to vouch for UC" }K<sup>-1</sup>Gavin
  - { "Carol Christ's public key is 0x... and I trust her to vouch for UCB" }K<sup>-1</sup>Mike
  - { "John Canny's public key is 0x... and I trust him to vouch for CS" }K<sup>-1</sup>Carol
  - { "Nick Weaver's public key is 0x..." }K<sup>-1</sup>John

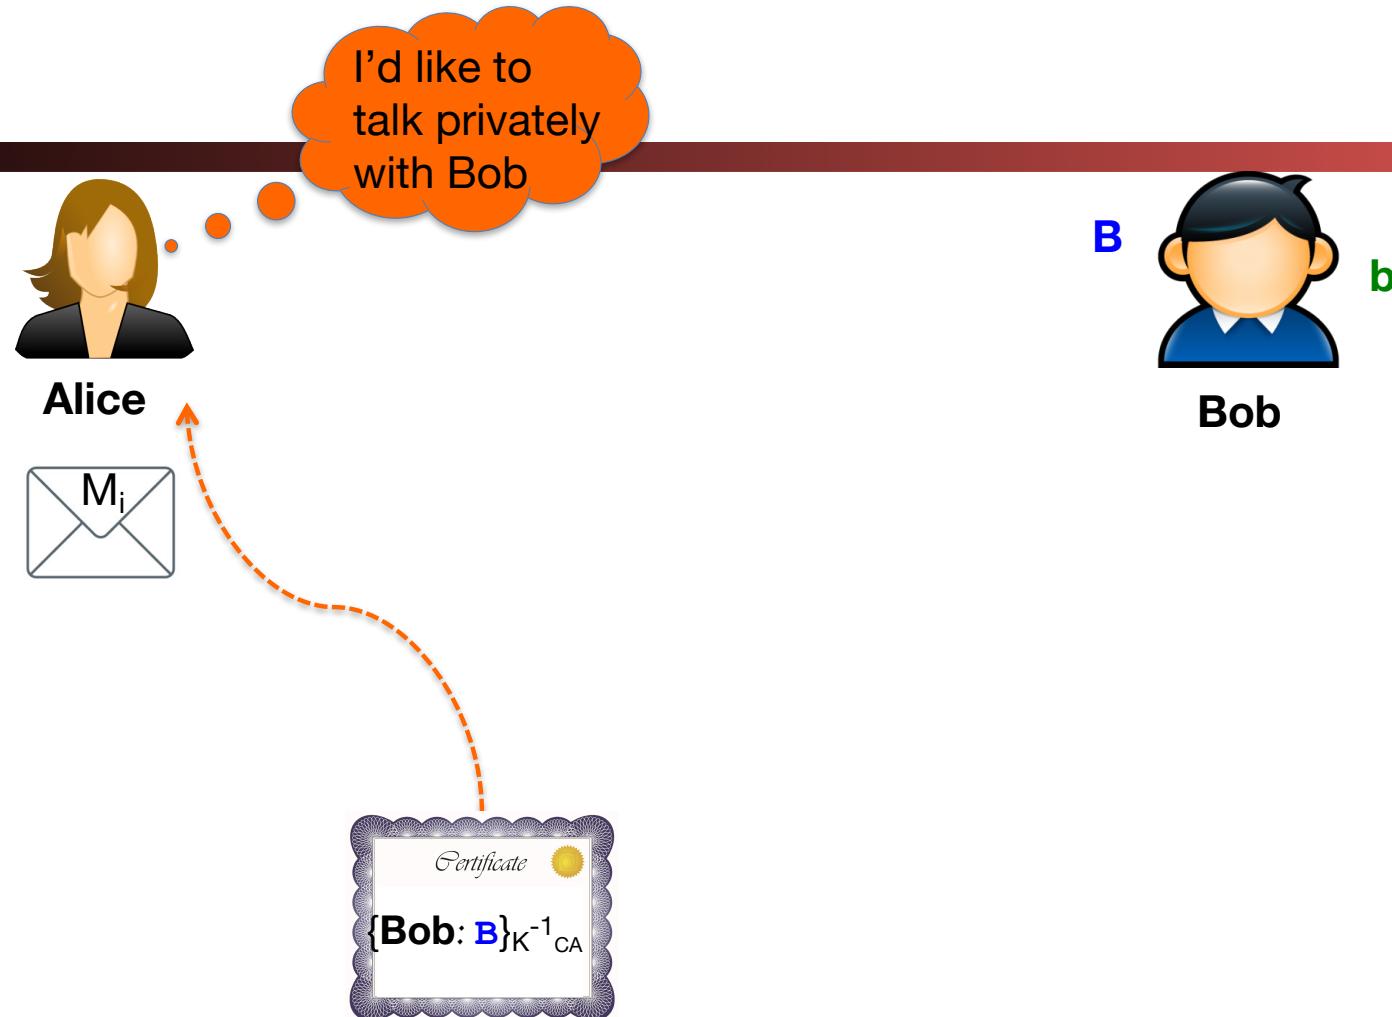
# Scaling Digital Certificates, con't

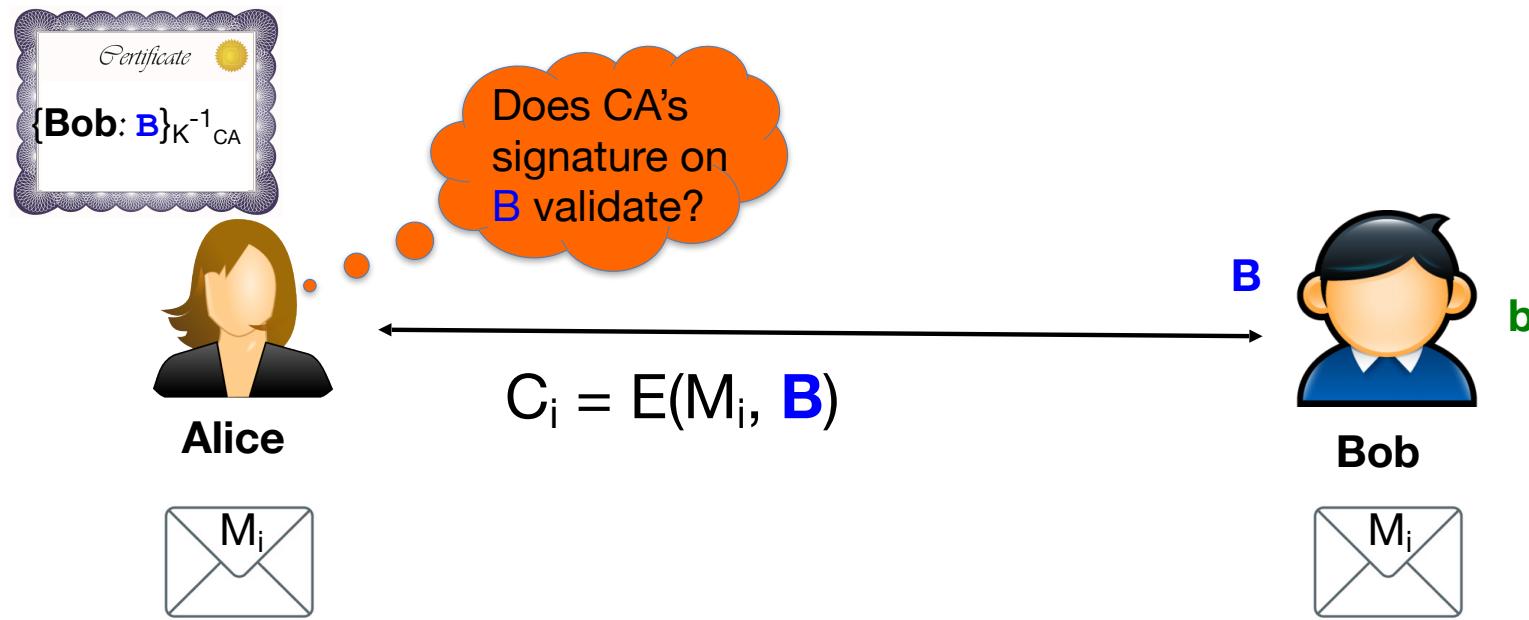
- I put this last certificate on my web page
  - (or shoves it under your door)
- Anyone who can gather the intermediary keys can validate the chain
  - They can get these (other than Gavin's) from anywhere because they can validate them, too
  - In fact, I may as well just include those certs as well, just to make sure you don't have to go search for them
- Approach #2: have multiple trusted parties who are in the business of signing certs ...
  - (The certs might also be hierarchical, per Approach #1)

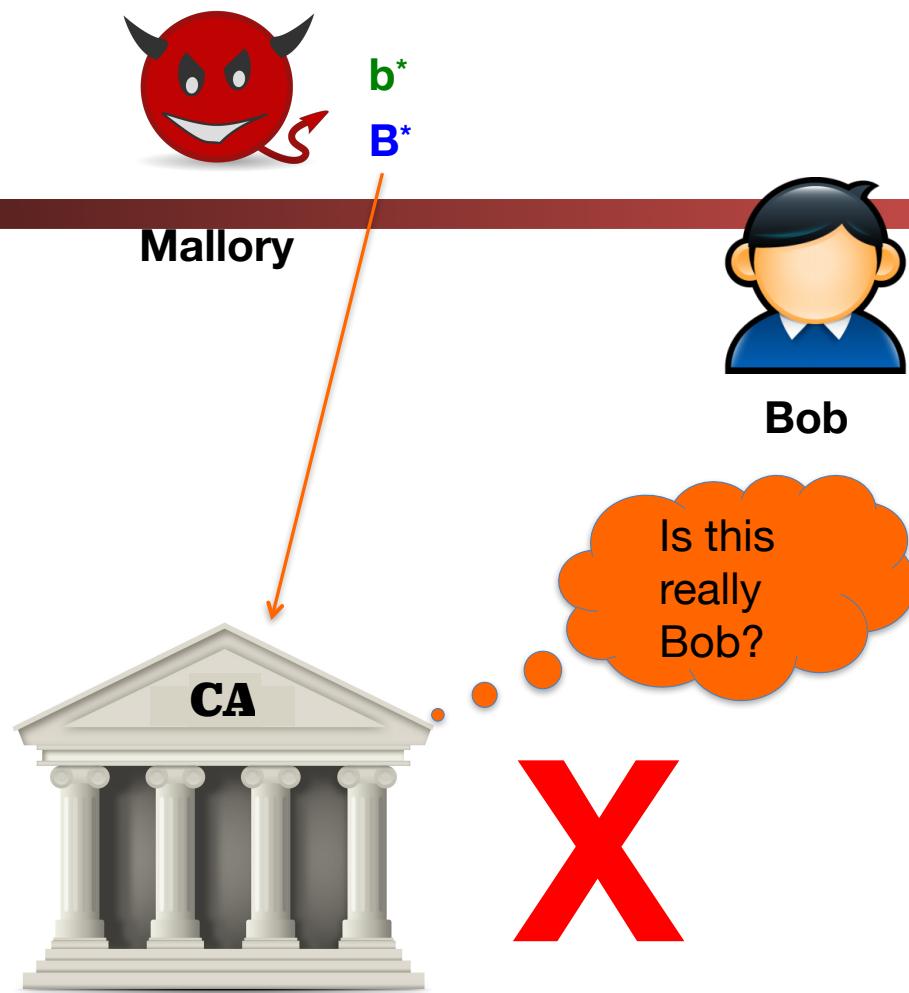
# Certificate Authorities

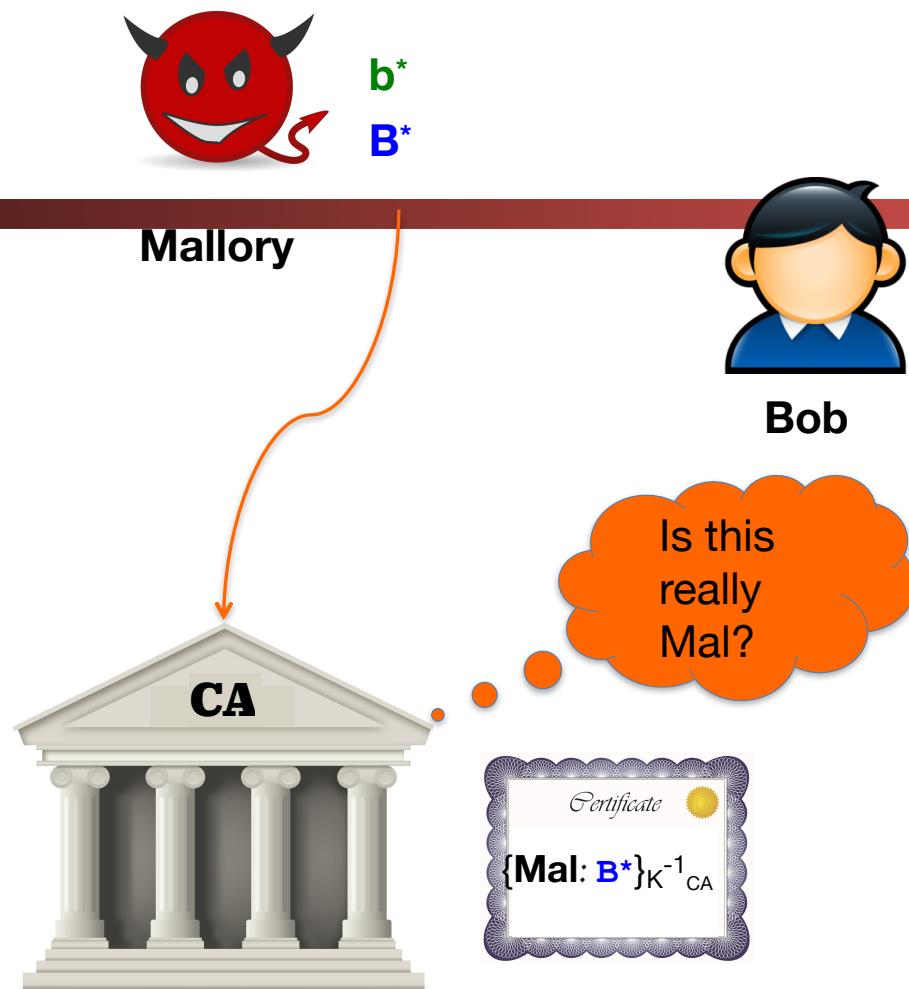
- CAs are trusted parties in a Public Key Infrastructure (PKI)
- They can operate offline
  - They sign (“cut”) certs when convenient, not on-the-fly (... though see below ...)
- Suppose Alice wants to communicate confidentially w/ Bob:
  - Bob gets a CA to issue {Bob’s public key is B}  $K^{-1}_{CA}$
  - Alice gets Bob’s cert any old way
  - Alice uses her known value of  $K_{CA}$  to verify cert’s signature
  - Alice extracts B, sends  $\{M\}K_B$  to Bob

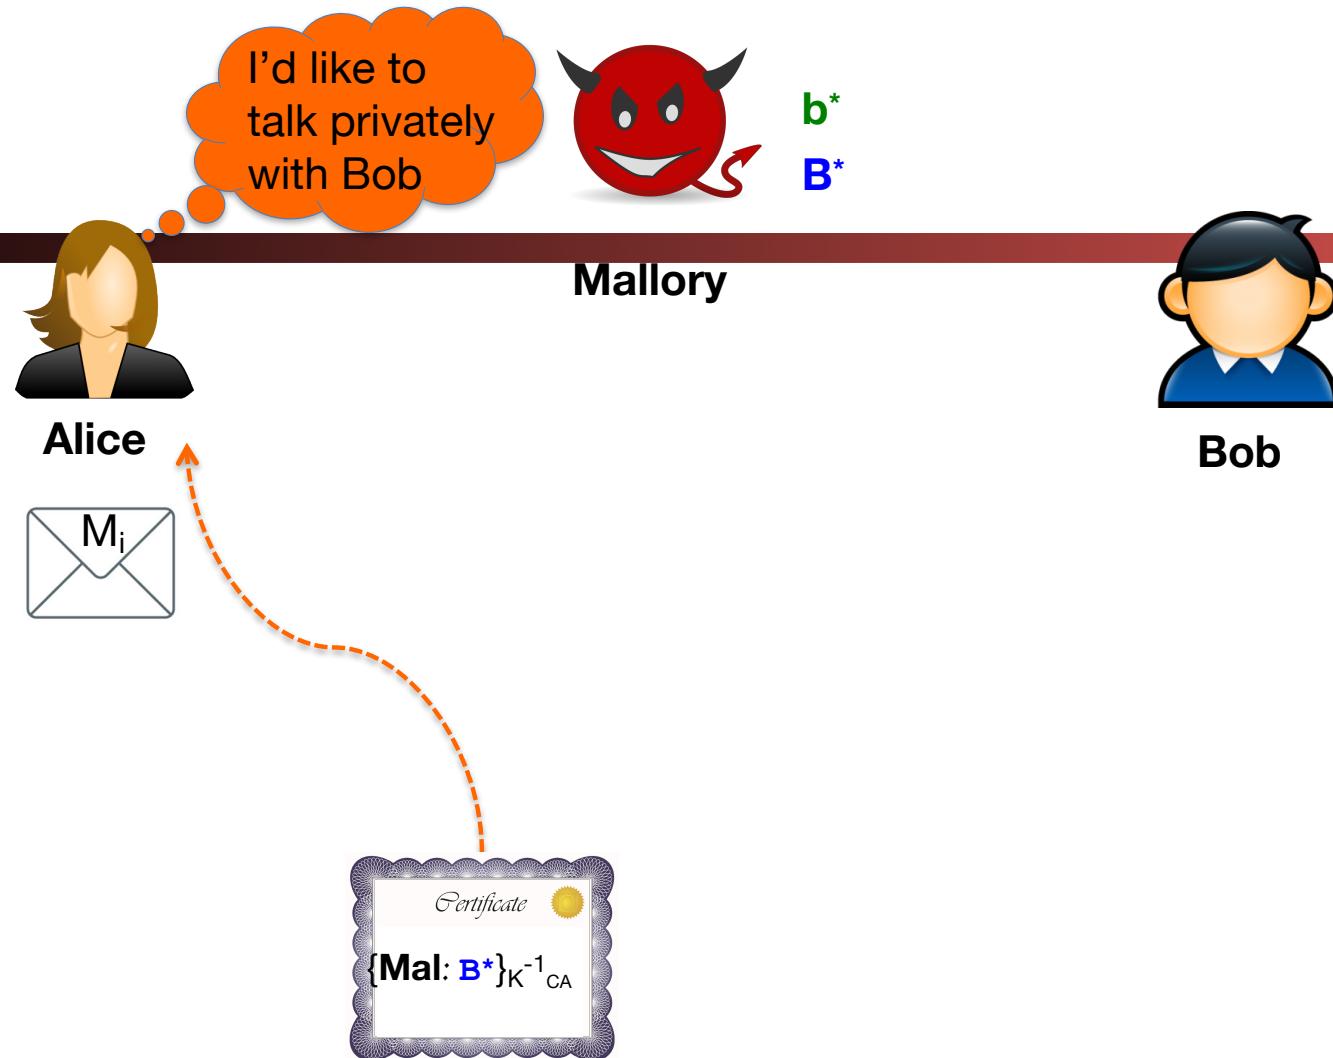














Alice



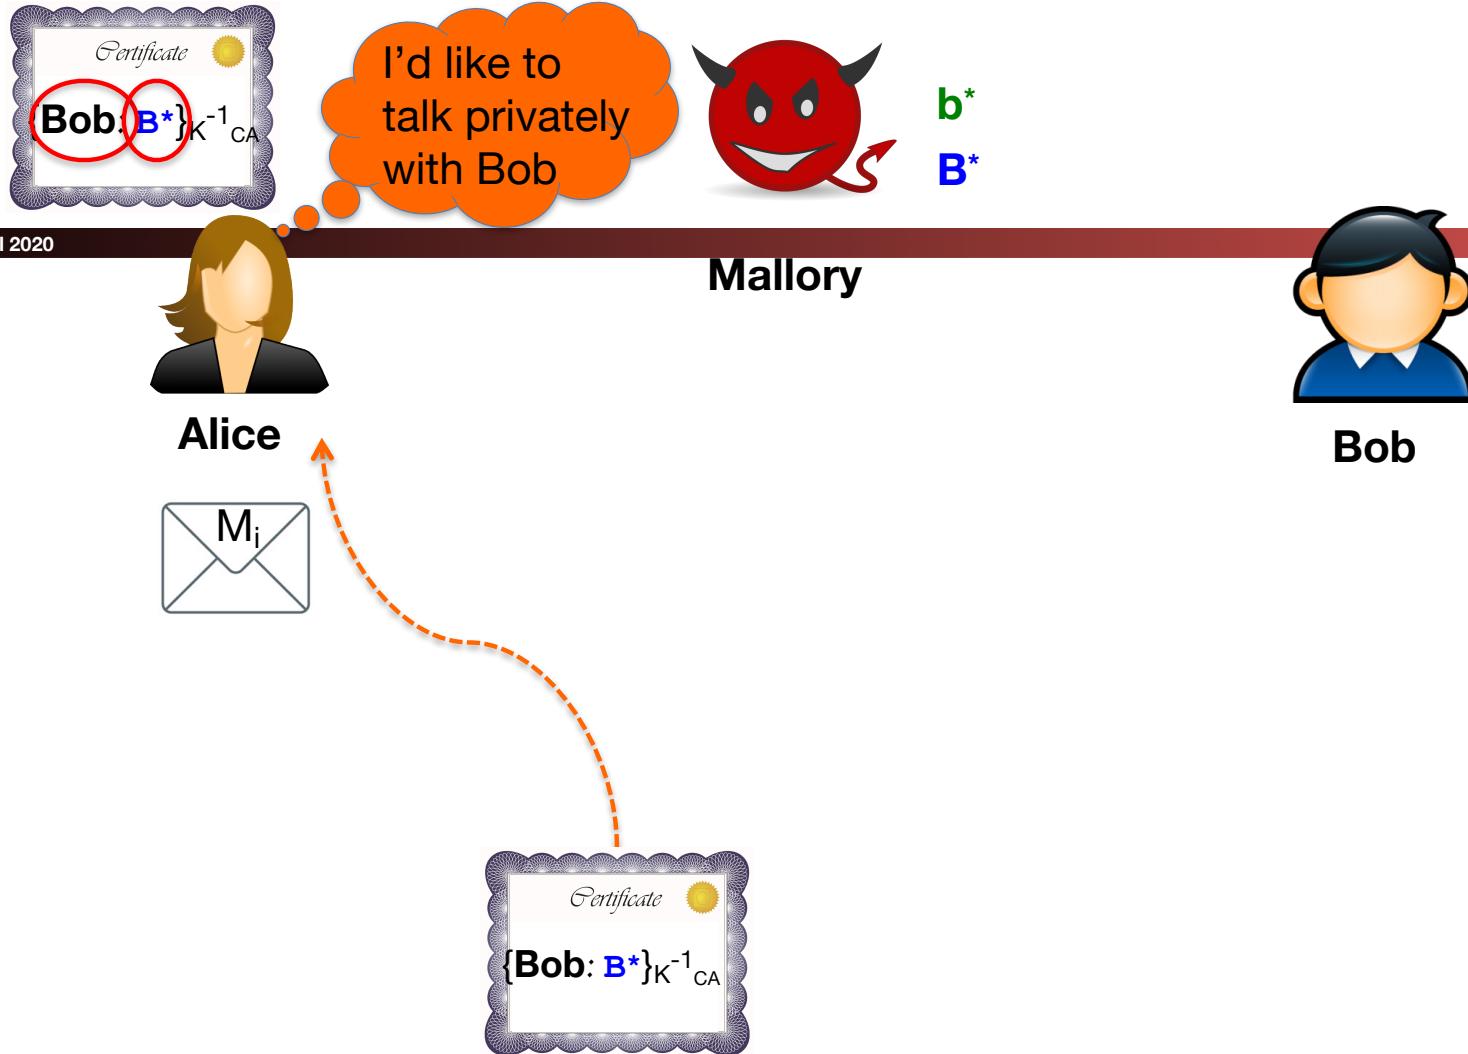
Mallory

 $b^*$  $B^*$ 

Bob

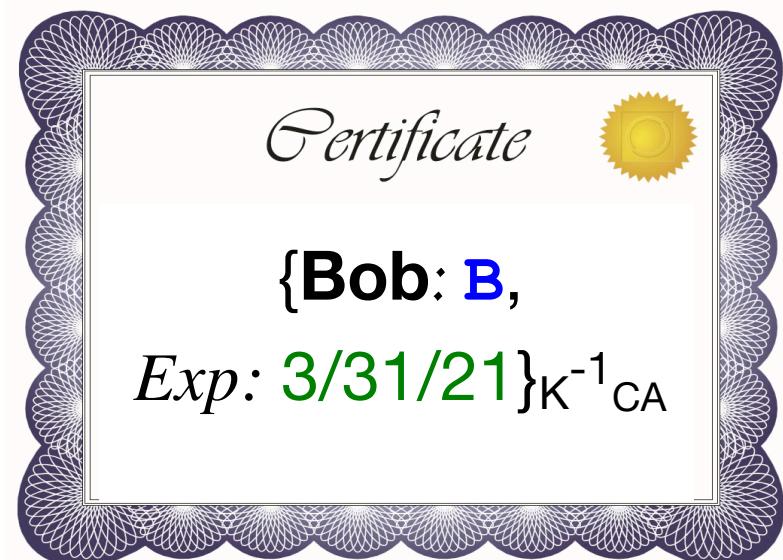
# Revocation

- What do we do if a CA screws up and issues a cert in Bob's name to Mallory?



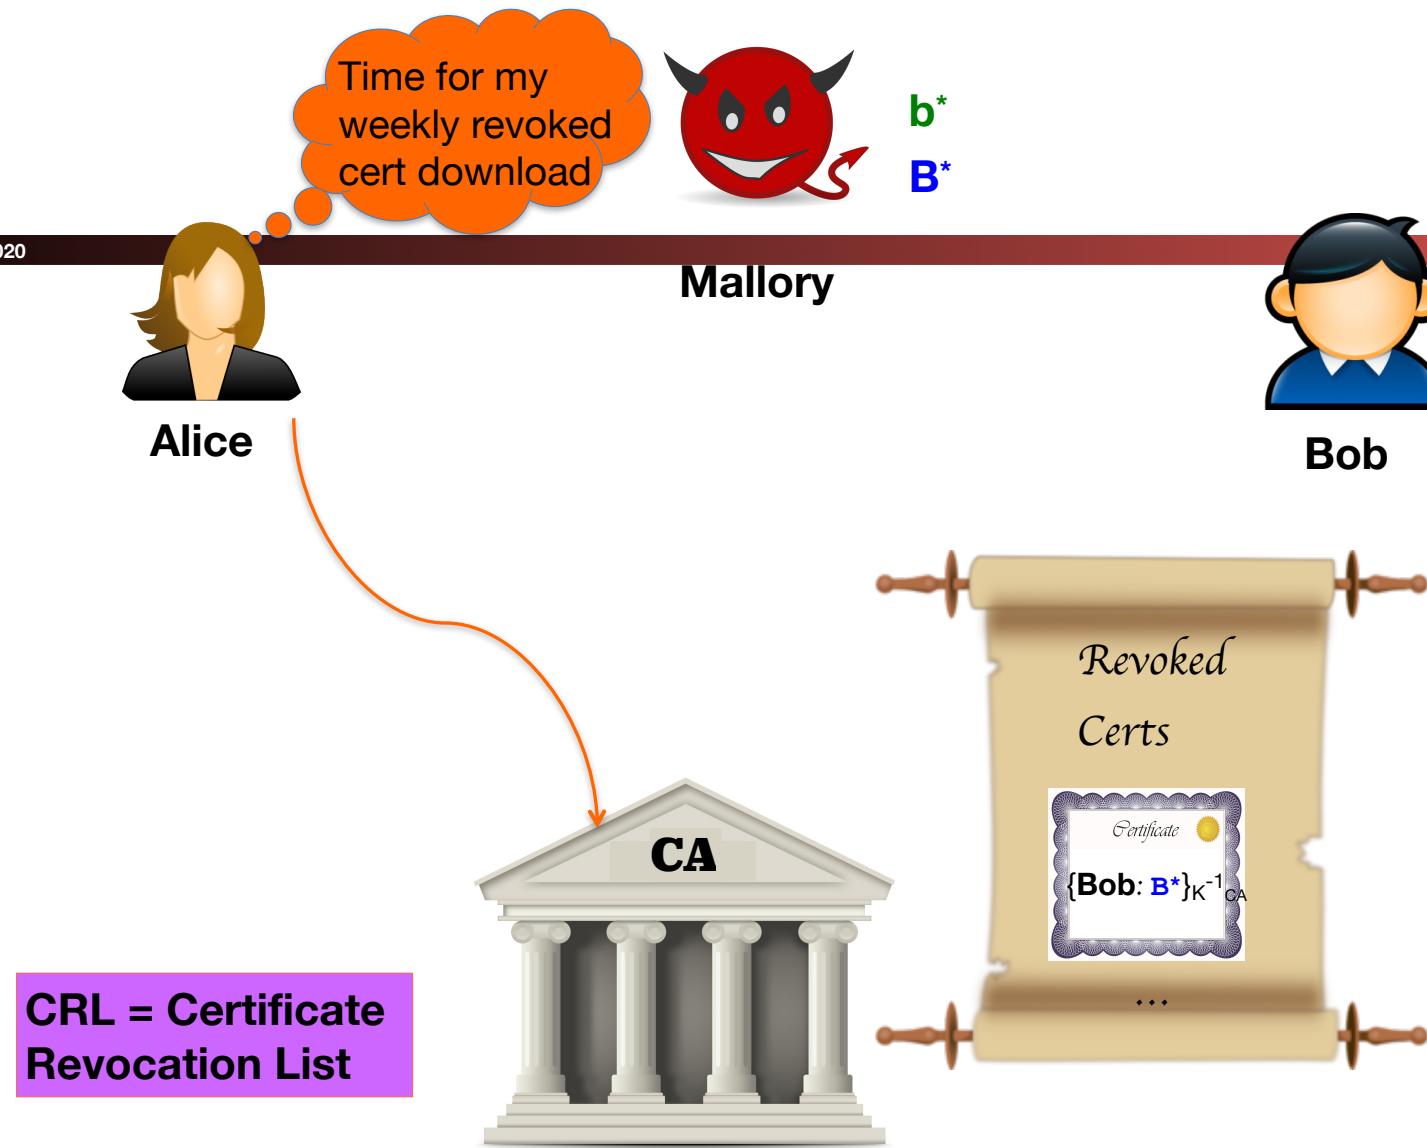
# Revocation

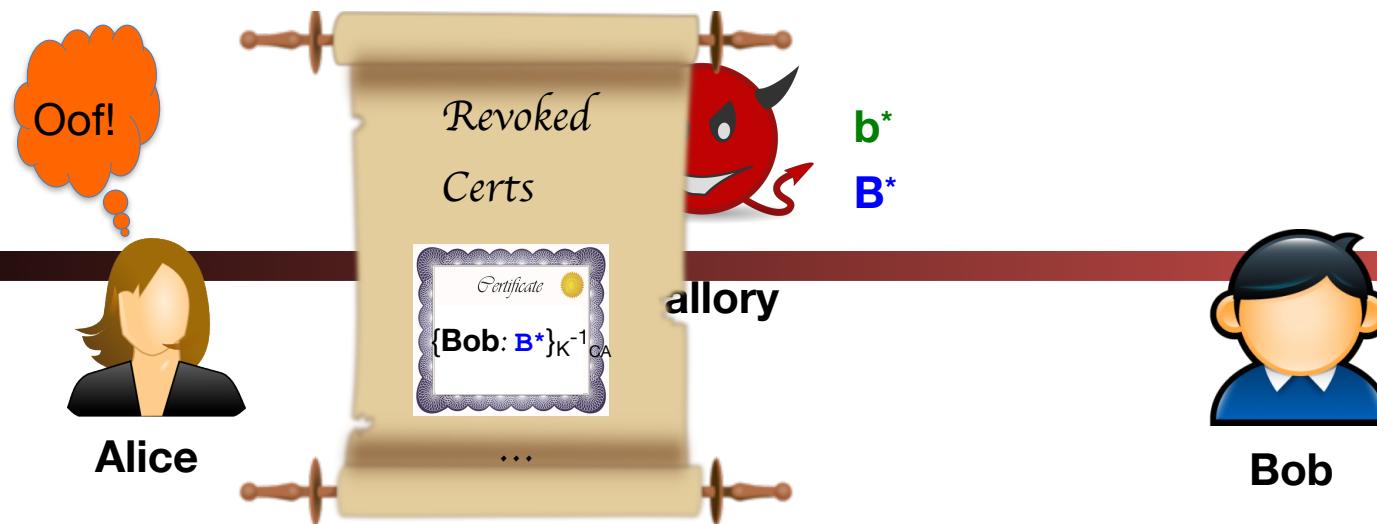
- What do we do if a CA **screws up** and issues a cert in Bob's name to Mallory?
  - E.g. Verisign issued a **Microsoft.com** cert to a **Random Joe**
  - (Related problem: Bob realizes **b** has been **stolen**)
- *How do we recover from the error?*
- **Approach #1: expiration dates**
  - Mitigates possible damage
  - But adds management burden
    - Benign failures to renew will break normal operation
    - LetsEncrypt decided to make this VERY short to force continual updating



# Revocation, con't

- Approach #2: announce revoked certs
  - Users periodically download cert revocation list (CRL)



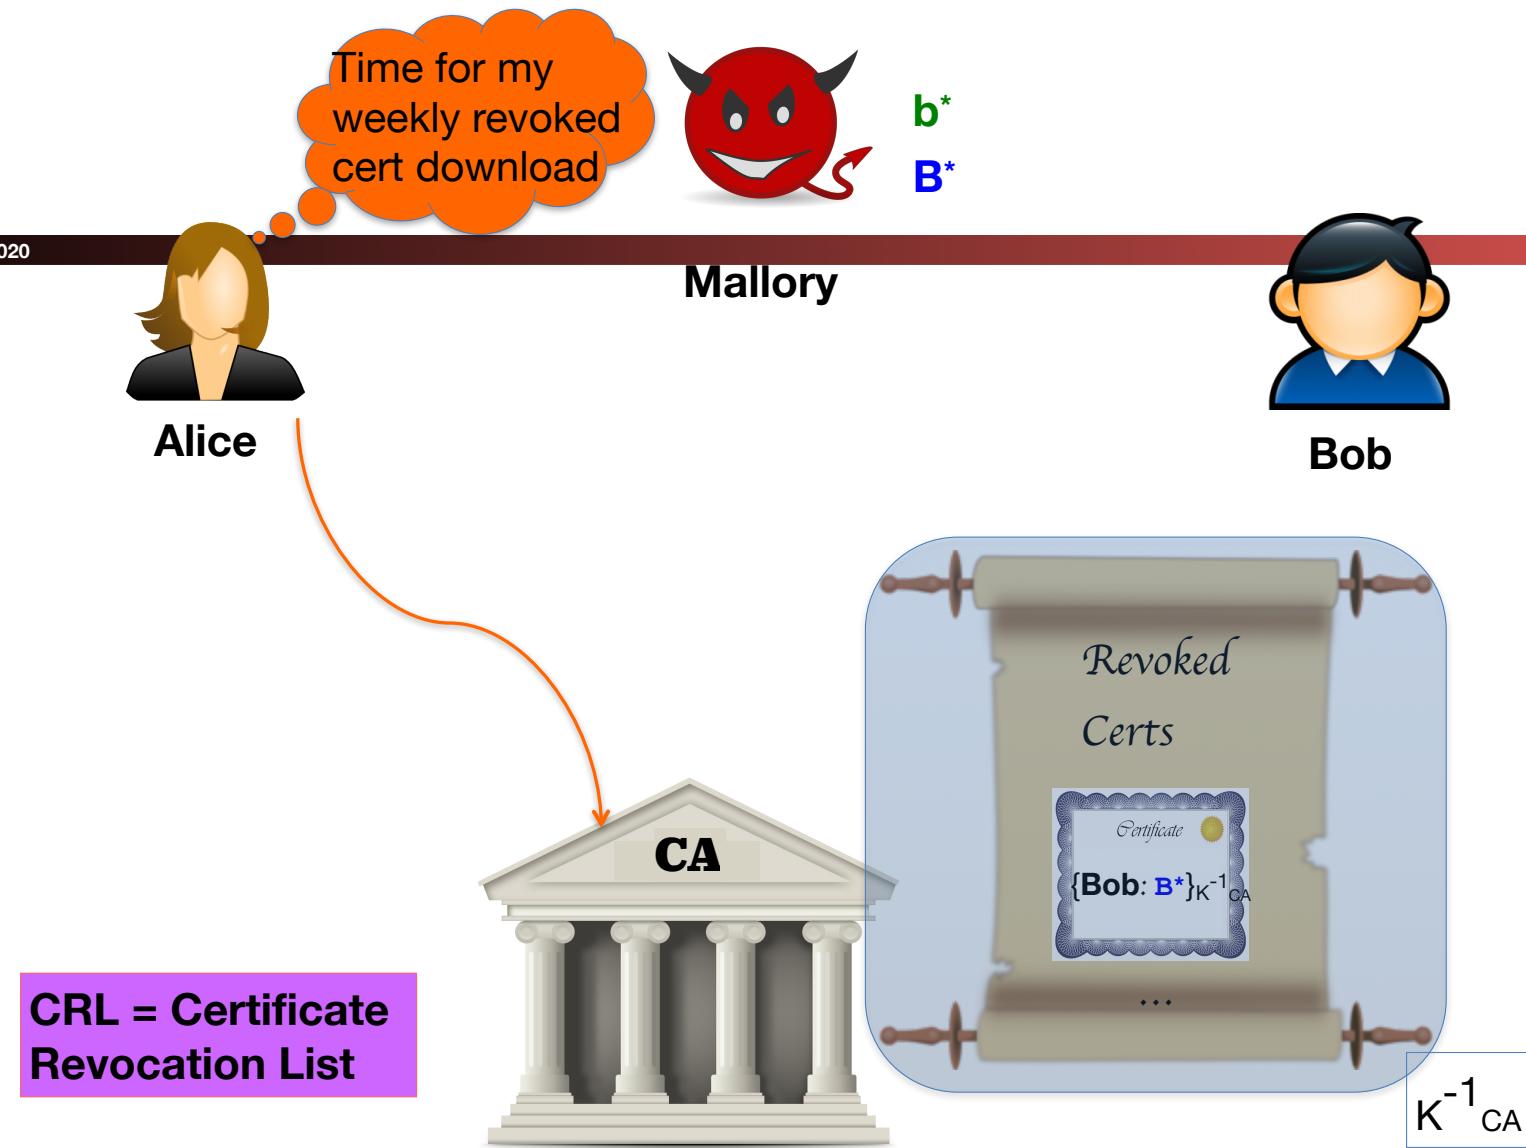


CRL = Certificate Revocation List



# Revocation, con't

- Approach #2: announce revoked certs
  - Users periodically download cert revocation list (CRL)
- Issues?
  - Lists can get large
  - Need to authenticate the list itself – how?



# Revocation, con't

- Approach #2: announce revoked certs
  - Users periodically download cert revocation list (CRL)
- Issues?
  - Lists can get large
  - Need to authenticate the list itself – how? Sign it!
  - Mallory can exploit download lag
  - What does Alice do if can't reach CA for download?
    - Assume all certs are invalid (fail-safe defaults)
      - Wow, what an unhappy failure mode!
    - Use old list: widens exploitation window if Mallory can “DoS” CA (DoS = denial-of-service)



# Biggest Problem is Often Complexity

- The X509 "standard" for certificates is incredibly complicated
  - Why? Because it tried to do everything...
  - If you want your eyes to bleed...
  - <https://tools.ietf.org/html/rfc5280>
  -

# The (Failed) Alternative: The “Web Of Trust”

- Alice signs Bob's Key
  - Bob Signs Carol's
- So now if Dave has Alice's key, Dave can believe Bob's key and Carol's key...
  - Eventually you get a graph/web of trust...
- PGP started out with this model
  - You would even have PGP key signing parties
  - But it proved to be a disaster:  
Trusting central authorities can make these problems so much simpler!